PREDICTING YOUTH SUBSTANCE USE USING DECISION TREES

DATA 5322- Statistical Machine Learning -2

- By Sravya Murala

INTRODUCTION

Data: NSDUH 2023

- National Survey on Drug Use and Health
- 10,561 respondents under age 18

The three questions being answered are:

- Did this youth ever use alcohol? (Binary)
- How frequently do marijuana users consume it? (Multi-class)
- How many days per year is marijuana used? (Regression)

THEORETICAL BACKGROUND (DECISION TREES)

- Recursively splits data by the feature that best reduces impurity
- Classification uses Gini index, Regression uses RSS
- Results in an interpretable flowchart of decisions

Tuning: Tree depth is controlled using cross-validation pruning, cost-complexity parameter aplha

Evaluation: Accuracy / confusion matrix for classification, RMSE for regression

Limitation: Single trees are high variance and prone to overfitting

THEORETICAL BACKGROUND (ENSEMBLES)

Model	How it works	Key Parameter	Limitation
Bagging	Bootstrap + average many trees	mtry = all features (p)	Trees are correlated
Random Forest	Bagging + random feature subset per split	mtry = √p	Less interpretable
Gradient Boosting	Sequentially corrects errors of prior trees	n.trees, shrinkage, depth	Sensitive to tuning

METHODOLOGY: DATA PREPARATION

Variable Selection:

- 16 predictors chosen from peer, parental, school, and demographic categories
- Avoided leakage an example is did not use "age first used alcohol" to predict "ever used alcohol"

Cleaning Steps:

- Removed special coded non-answers (991 = never used, 993 = not used recently)
- Converted all categorical variables to factors
- Multi-class filtered to marijuana users only (MRJFLAG = 1)

Class Imbalance:

- Binary outcome: (77% No, 23% Yes) addressed using ROSE oversampling
- Balanced training set: 50/50 split

Train/Test Split: 80% training, 20% testing across all three

HYPERPARAMETER TUNING

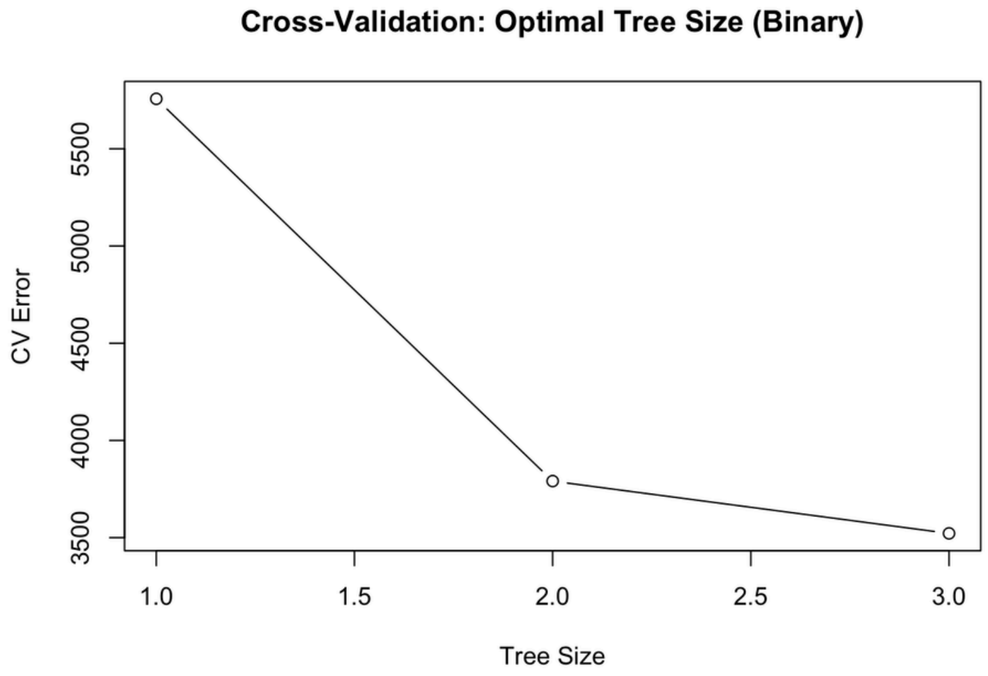

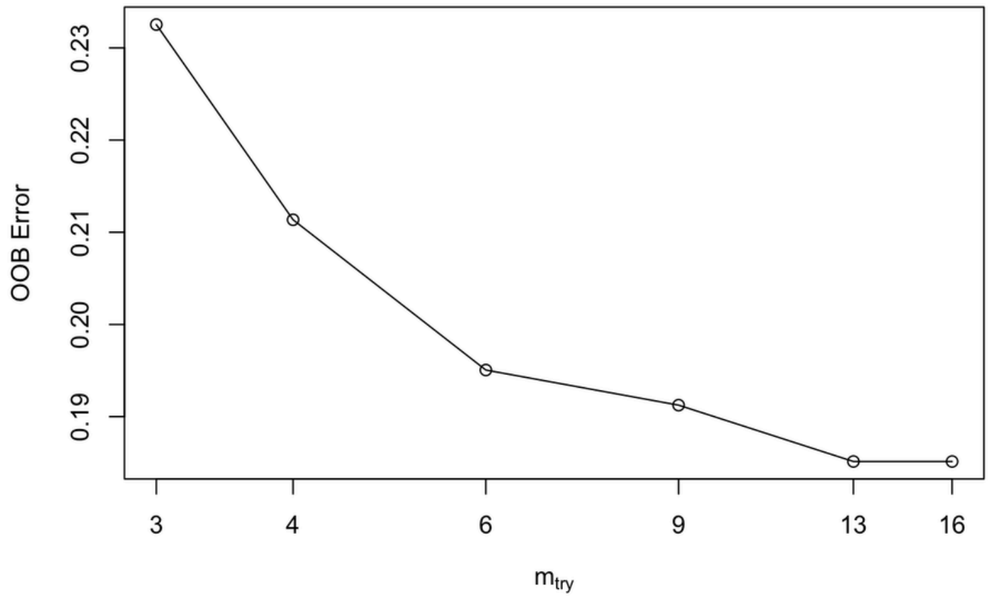

HYPERPARAMETER TUNING

Decision Tree Pruning:

- Used cross-validation misclassification error
- Binary: optimal size = 3 nodes
- Regression: optimal size = 3 nodes

Random Forest mtry:

- Used tuneRF with stepFactor = 1.5
- Searched mtry values: 3, 4, 6, 9, 13, 16
- OOB error decreased with higher mtry

Gradient Boosting:

- Fixed: n.trees = 500, shrinkage = 0.01, interaction.depth = 3
- Best number of trees selected via 5-fold cross-validation

COMPARISION

Encoding Type	Variable	What it measures	Best for
Binary (0/1)	ALCFLAG	Ever used alcohol or not	Simple yes/no classification
Ordinal factor	MRJYDAYS (1–6)	Ordered frequency categories	Multi-class when order matters
Continuous integer	IRMJFY (1–365)	Exact days of marijuana use per year	Regression when precision needed

TREE INTERPRETATION

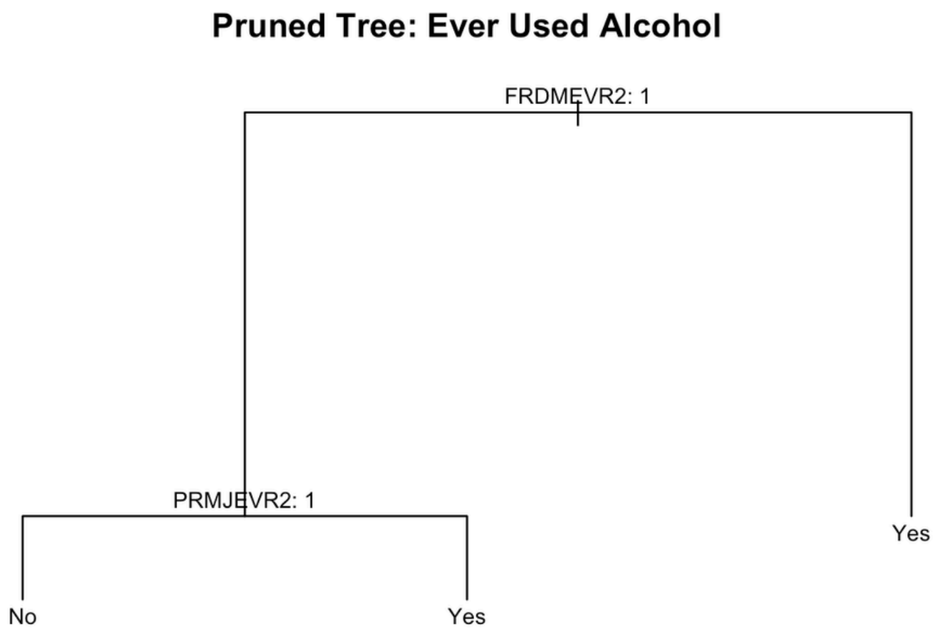

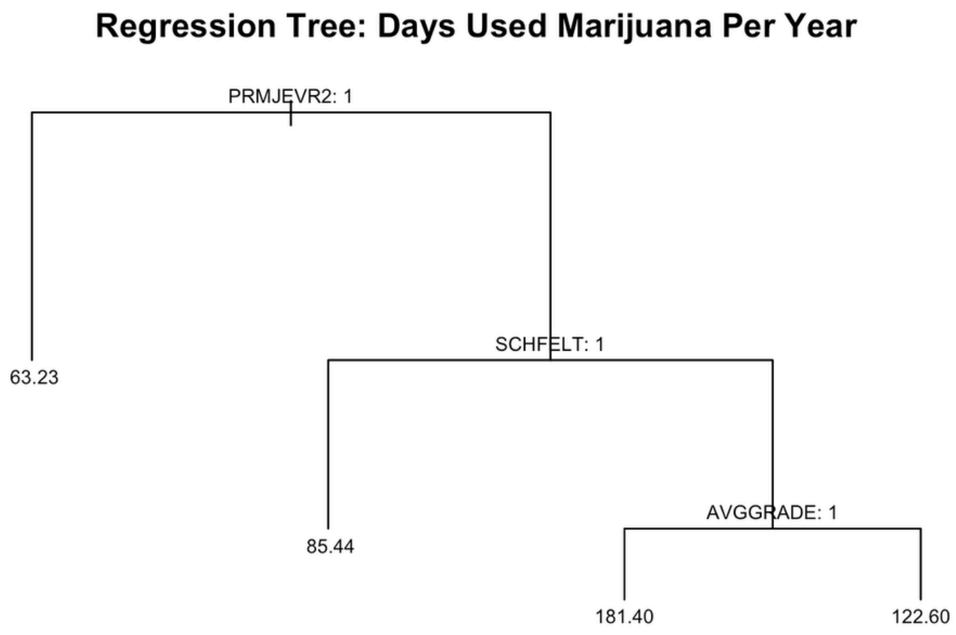

FRDMEVR2 = Friends disapprove of marijuana use
PRMJEVR2 = Parent told youth marijuana is wrong
SCHFELT = How youth feels about school
AVGGRADE = Average grade in school

RESULTS: BINARY CLASS

Model	Accuracy
Basic simple Tree (imbalanced)	77.20%
Balanced Tree	73.60%
Pruned Tree	73.60%
Bagging	70.80%
Random Forest	**73.90%**
Boosting	71.90%

RESULTS: MULTI-CLASS AND REGRESSION

Multi-Class	
Model	**Accuracy**
Basic Tree	23.30%
Pruned Tree	23.30%
Bagging	25.30%
Random Forest	24.70%
Boosting	**29.10%**

Regression	
Model	**RMSE**
Basic Tree	**109.6**
Pruned Tree	110.6
Boosting	110.4
Random Forest	114.8
Bagging	117.8

VARIABLE IMPORTANCE

Variable Importance of Alcohol

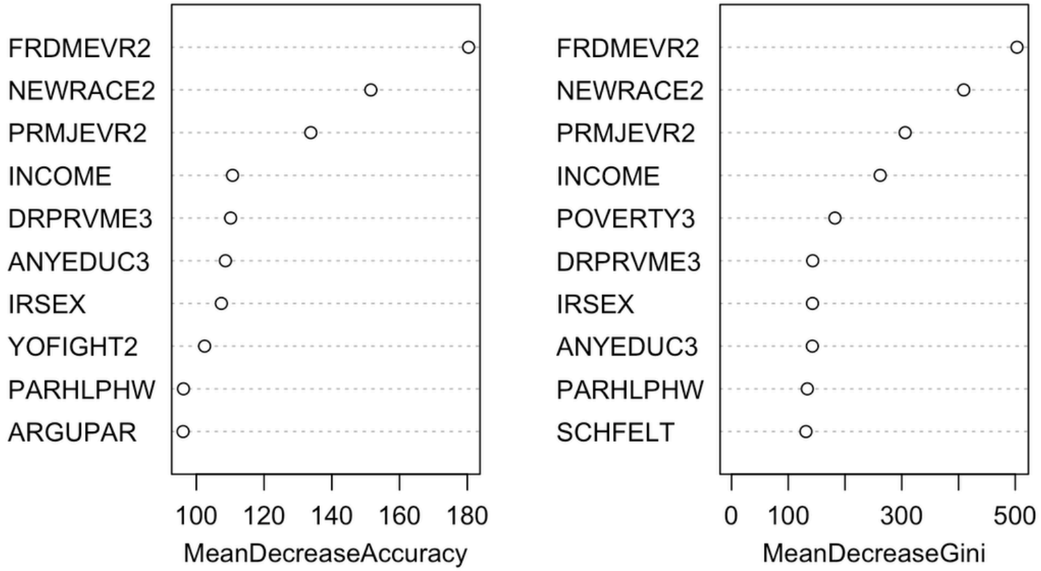

Variable Importance of Marijuana Frequency

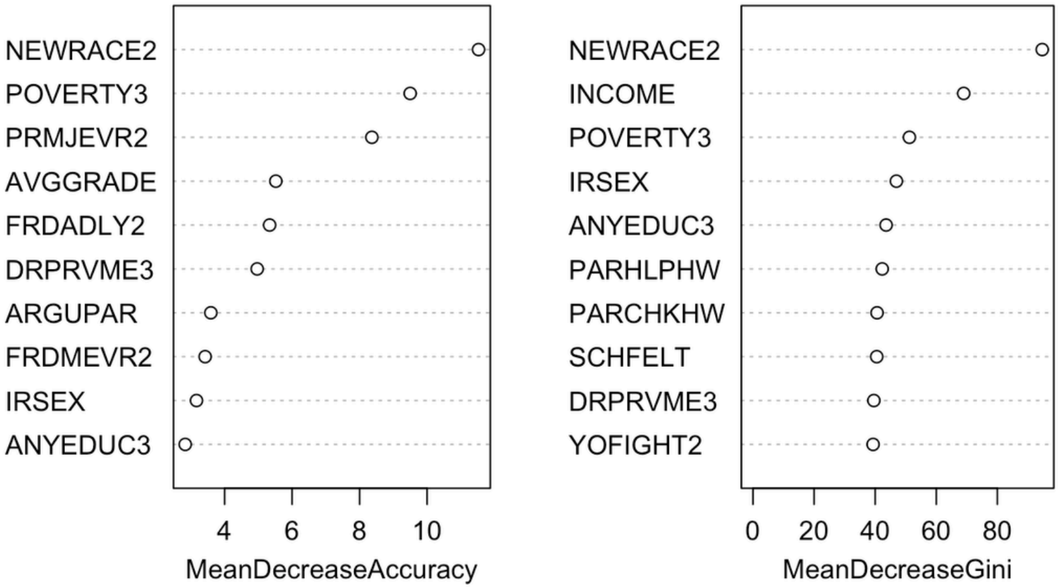

VARIABLE IMPORTANCE

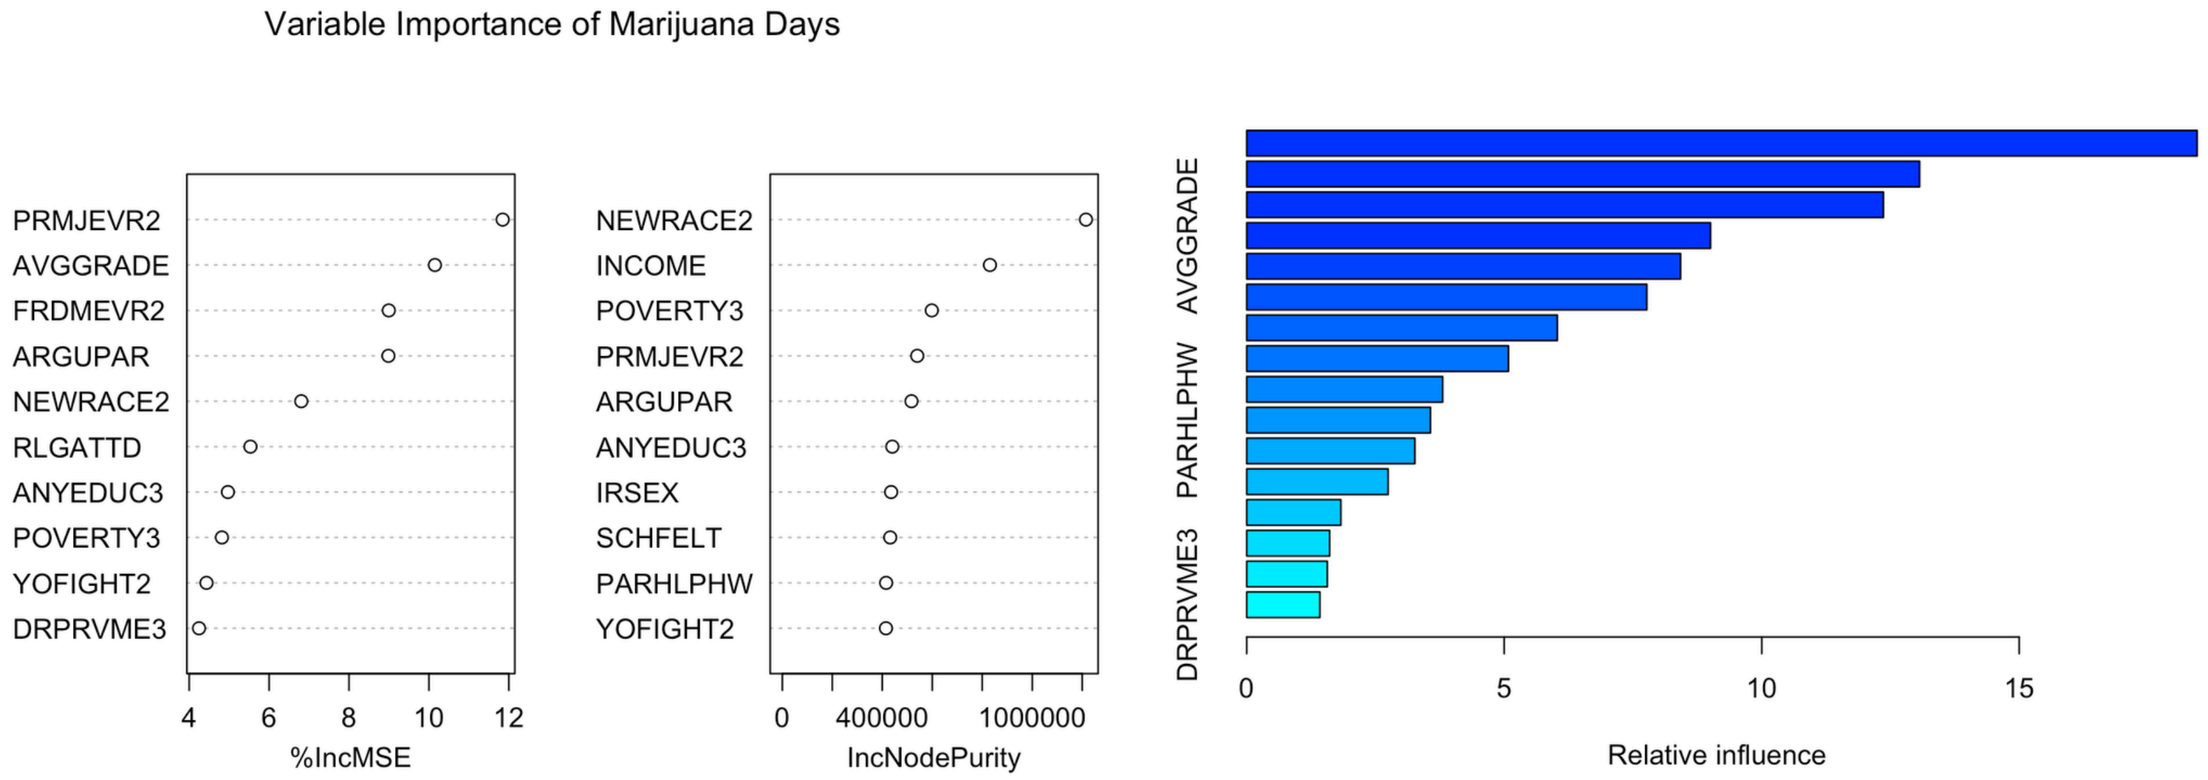

DISCUSSION

Performance of Model:
Random Forest most effective for binary (73.9%), Boosting most effective for multi-class (29.1%)
Basic tree proved quite competitive for the regression problem but not always superior

Reasons behind poor performance of model:
Multiclass classification: 6 overlapping frequency bins difficult to classify based on survey responses
Regression problem: Survey respondents can report days between 1-365

Limitations:
Response bias in the survey data
Binary nature of most features, low granularity
Correlation but no causation
Race and income as features indicative of systemic issues

Future Work:
Include psychological feature
SMOTE technique for balance over ROSE

CONCLUSION

- Family and peer influences have greater predictive power than mere demographics
- Simple trees had grasped the important patterns, ensemble methods helped with stability, but not much more
- The prediction of frequency, the multi-class is much more difficult than a binary classification
- Drug use behavior is very complicated, and this is just one step toward understanding the problem
- Implication: Preventive measures that emphasize family and peer environments may prove more successful than mere demographics

REFERENCES

1. Substance Abuse and Mental Health Services Administration. National Survey on Drug Use and Health 2023. U.S. Department of Health and Human Services, 2023.
2. James, G., Witten, D., Hastie, T., & Tibshirani, R. An Introduction to Statistical Learning. Springer, 2023.
3. Liaw, A. & Wiener, M. (2002). Classification and Regression by randomForest. R News, 2(3), 18–22.
4. Greenwell, B. et al. (2020). gbm: Generalized Boosted Regression Models. R package version 2.1.8.
5. Menardi, G. & Torelli, N. (2014). Training and assessing classification rules with imbalanced data. Data Mining and Knowledge Discovery, 28, 92–122.

THANK YOU